

# Proactive Computer Security

## Shellcode

Ken Friis Larsen  
<kflarsen@diku.dk>

Department of Computer Science  
University of Copenhagen

May 2, 2016

Original slides by Morten Brøns-Pedersen <f@ntast.dk>  
and Morten Shearman Kirkegaard <moki@fabletech.com>

## Disclaimer

These slides do not replace the articles, and you should read those as well. The slides are primarily an aid for me to get my points across during the lecture. Don't hesitate to ask questions during or after the lecture, or on the forum.

# Today's Program

- ▶ Shellcode vs. programs
- ▶ x86 assembler code
- ▶ Considerations when writing shellcode

The goal of an attacker is to take control of some privileged process, so he can act with the privileges of that process.

- ▶ Setuid binary (e.g., backup software) – other user / root
- ▶ Local daemon (e.g., cron) – other user / root
- ▶ Remote server (e.g., web server) – other host
- ▶ Remote client (e.g., web browser) – other host

The ultimate control of a target process is known as *arbitrary code execution*

In that case, an attacker can choose a sequence of instructions to be executed in the target process. The instructions are injected into the process as part of the exploitation. Traditionally the code has spawned a shell.

# Question

How does a “sequence of instructions” (i.e., shellcode) differ from an executable file (e.g., `C:\windows\system32\cmd.exe` or `/bin/sh`)?

A program is not just a sequence of instructions. Apart from code most programs also need static and dynamic data, some way to call library functions, and so on.

Both Windows and UNIX have several executable formats, but the most used ones are:

- UNIX ELF (Executable and Linkable Format)

- Windows PE (Portable Executable)

Shellcode can be more advanced than to spawn a shell, but it is still referred to as “shellcode.”

Common shellcode categories:

- ▶ Execve(/bin/sh) – only local
- ▶ Listen Shell
- ▶ Connect-back Shell
- ▶ Downloader
- ▶ Egg Hunter
- ▶ Socket Hunter



## Typical shellcode

- ▶ Is tiny — often less than 100 bytes
- ▶ Contains no NUL bytes
- ▶ Is position independent
- ▶ Does not rely on a specific address space layout

To write good shellcode you need to know your processor and your operating system. Source code is great. Reverse engineering is good enough.

We will look at Intel 80386 shellcode with focus on Linux. The same basic principles apply to other processors and operating systems.

# 80386 Crash Course

Intel introduced the 32 bit 80386 processor in 1985. Nothing has changed since then.

This is – technically – not accurate. Advanced Micro Devices introduced the 64 bit AMD64 processor series in 2003.

The architecture of the Intel 80386 processor is officially called *IA32*, but is often referred to as *x86* or *i386*.

The *AMD64* architecture is often referred to as *x86-64*, or – by Intel – as *IA-32e* or *EM64T* and – by Microsoft – as *x64*.

Do **not** confuse it with *IA64* (Itanium is an *IA64* processor).

## Intel x86 assembler.

- ▶ Before you do anything else, get a copy of “Intel Architecture Software Developer’s Manual”.
- ▶ Two syntaxes: AT&T and Intel (we use Intel).
- ▶ Little endian.

# Question

What does it mean that the processor is “little endian”?

# Question

Actually while we're talking about bits, what is "two's complement representation"? (And why is it smart?)

Extra credit: What 8-bit number is this:

$$11110100 = -12 \text{ or } 244$$



There are 8 (Almost) general purpose 32bit registers.  
See also “The Art of Picking Intel Registers”<sup>1</sup>.

**EAX** Accumulator. Results.

**ECX** Loop counter.

**EDX** Data register, I/O pointer. 64-bit extension of EAX.

**EBX** Base register (data pointer). General purpose.

**ESP** Stack pointer.

**EBP** Base pointer.

**ESI** Source register.

**EDI** Destination register.

---

<sup>1</sup><http://www.swansontec.com/sregisters.html>.

For each register there are one or more extra “registers” which are a subset of the full register. They are:

Register	Lower 16 bits	Lower 8 bits	bits 8–15
EAX	AX	AL	AH
ECX	CX	CL	CH
EDX	DX	DL	DH
EBX	BX	BL	BH
ESP	SP	–	–
EBP	BP	–	–
ESI	SI	–	–
EDI	DI	–	–

Other registers:


**EIP** Instruction pointer. Can't be controlled directly.

**Eflags** Indicate events. For example ZF (zero flag), SF (sign flag) and CF (carry flag). See “Intel Architecture Software Developer's Manual, Volume 1”.


When the 80386 is running in *Protected Mode* all memory references are 48 bit; a 16 bit segment selector, and a 32 bit offset.

The 16 bit segment selector is loaded into a segment register.

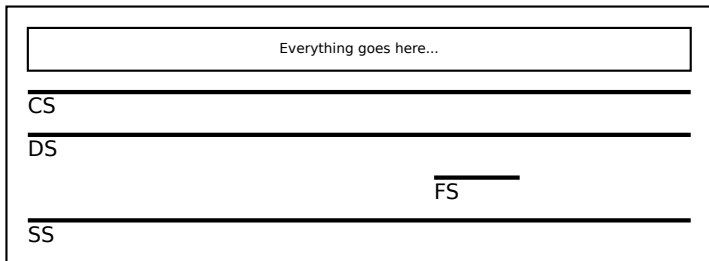
- ▶ Instructions are read from the segment in **CS**
- ▶ The stack is accessed via the segment in **SS**
- ▶ Data is, by default, accessed via the segment in **DS**
- ▶ Data can be accessed via the segments in **ES**, **FS** and **GS** by using an instruction prefix.



## What Intel Imagined:



## What Operating System Designers Did:



Segmentation is practically disabled, CS:12345678 is DS:12345678.  
Code is data and data is code. Pointers are 32 bits.

One exception is a segment register used for TLS. **FS:0000**1234  
might be **DS:AAAA**1234 in thread A and **DSBBBB**1234 in thread B.



## Execution:

1. The instruction at CS:EIP is fetched from memory
2. EIP is incremented by the length of the instruction
3. The instruction is executed
4. Rinse, repeat

Some instructions – `call` and `jmp` – change EIP (either by offset or by absolute address). This is done *after* adding the length of the instruction to EIP.

- ▶ The instruction at CS:EIP is fetched from memory
- ▶ EIP is incremented by the length of the instruction
- ▶ `call` pushes EIP (the *return address*) onto the stack
- ▶ The offset is added to EIP

# Question

What does these instructions do?

EB 00    `jmp short 0`    is a no-op.

EB FE    `jmp short -2`    is an infinite loop.

## Some OP codes

**PUSH src** Pushes its operand onto the stack. Decreases ESP according to the operand. If ESP itself is pushed, it is the value of ESP *before* the operation that gets pushed<sup>2</sup>.

**POP dst** Pops a value from the stack and stores it in the operand. Increases ESP accordingly.

**MOV dst, src** Moves the value of *src* to *dst*. The source operand can be a value, a register or a memory location (using the  $[reg1 + reg2 \cdot imm1 + imm2]$  notation). The destination can be a register or a memory location.

**ADD dst, src** Adds *src* and *dst* and stores the result in *dst*.

---

<sup>2</sup>GOTCHA: **PUSH BYTE *imm*** pushes 4 bytes (with the first byte sign-extended into the upper three), but **PUSH WORD *imm*** pushes only 2 bytes.

# Stack?

What is “the stack” and what is it used for? (We saw one example already.)

## More OP codes

**SUB** *dst, src* You guessed it.

**CMP** *op1, op2* Compares *op1* to *op2* by subtracting *op2* from *op1*, and sets CF, OF, SF, ZF, AF and PF in Eflags accordingly. Normally used in conjunction with jump instructions.

**JZ, JNZ, JL, JLE, JG, JGE** Jump according to Eflags. If Eflags was set by **CMP** *op1 op2*, then the OP codes mean jump if *op1* is equal to, not equal to, less than, less than or equal to, greater than or greater than or equal to *op2* respectively. There are many other OP codes in the Jcc family.

**JMP** *dst* Jump to *dst* (register, memory, offset or absolute address).

## Even more OP codes

**CALL** *dst* Push the address of the following OP code on the stack and jump to *dst*.

**RET** Pop the return address and jump to it.

Look these up yourself:

LEA, PUSHA, POPA, PUSHF, POPF, CDQ, AND, OR, NOT, XOR, BSF, BSR, BSWAP, XCHG, MUL, DIV, NEG, NOP, INC, DEC, ENTER, LEAVE, LOOP, STOS, ROL, ROR, SHL, SHR, XADD, XLAT, and many more...


As we said before, get a copy of “Intel Architecture Software Developer’s Manual”, specifically “Volume 2: Instruction Set Reference”.

# System Calls



The shellcode must do something useful.

We need some help from the operating system.



```
AA000000 push byte 13 ; length
AA000002 push dword 0xAA110000 ; buffer
AA000007 push byte 1 ; fd
AA000009 call 0xAA220000 ; write()

AA110000 "hello, world\n"

AA220000 jmp [0xAA330000] ; Procedure Linkage Table
AA330000 0xBB000000 ; Global Offset Table

BB000000 mov edx, [esp+12] ; write() in libc
BB000004 mov ecx, [esp+8]
BB000008 mov ebx, [esp+4]
BB00000C mov eax, 4
BB000011 int 0x80
BB000013 ret
```

On Linux IA32 the system call number is passed to the kernel in EAX and the arguments in EBX, ECX, EDX, ESI, EDI and EBP. The result is returned in EAX.

The kernel system call handler is invoked via INT 0x80, SYSCALL or SYSENTER.

System call numbers don't change.

A list of the system calls can be found in  
`/usr/include/asm/unistd.h` (IA32) or  
`/usr/include/asm/unistd_{32,64}.h` (AMD64). (Paths might be  
system dependent)

...

<code>#define __NR_exit</code>	1
--------------------------------	---

<code>#define __NR_fork</code>	2
--------------------------------	---

<code>#define __NR_read</code>	3
--------------------------------	---

<code>#define __NR_write</code>	4
---------------------------------	---

<code>#define __NR_open</code>	5
--------------------------------	---

<code>#define __NR_close</code>	6
---------------------------------	---

...

```
write(STDOUT_FILENO, "hello, world\n", 13);
```

```
mov edx, 13          ; length
mov ecx, 0x11223344  ; buffer
mov ebx, 1           ; fd
mov eax, 4           ; SYS_WRITE
int 0x80
...
```

```
0x11223344:
```

```
db 'hello, world', 10
```

Use the man pages in section 2. They describe system calls.

man 2 exit

man 2 fork

man 2 read

man 2 write

...

man 2 execve

The following is true for Darwin, OpenBSD, FreeBSD, NetBSD, ...

On BSD IA32 the system call number is passed to the kernel in EAX and the arguments are passed on the stack, as if the kernel system call handler had been called via a regular call instruction; [ESP] is ignored, first argument is in [ESP+4], second in [ESP+8], and so on. The kernel system call handler is invoked via INT 0x80.



# Return Value

The result of a system call is returned in EAX, unless the call was to `SYS_FORK`.

The `fork()` library call returns

- ▶ 0 in the child process
- ▶ The PID of the child in the parent process

# Coffee Break

# Writing Shellcode

# Writing Position Independent Code.

Shellcode is often written to a buffer at some unknown address, and the dynamic linker won't help us locate anything.

The shellcode must either:

- ▶ Find its own absolute address
- ▶ Not rely on knowing the absolute addresses of itself

# Finding the Absolute Address.

It would be great to be able to do

```
EIP → ...  
EIP → mov eax, eip  
EAX EIP → ...  
EIP → add eax, 123  
EIP → ...  
      ...  
EAX → DATA
```

... but the 80386 has no such instruction.

## Using a trampoline

EIP → ...

EIP → jmp +123

EIP → ...

EIP → pop eax

EIP → ...

...

EIP → call -118

EAX → DATA

ESP →

...

ESP →

7A DA D5 AD

## Call without jump

EIP → ...

EIP → call +0

EAX EIP → ...

EIP → pop eax

EIP → ...

...

ESP →

...

ESP →

7A DA D5 AD

# Question

What are the pros and cons of a trampoline (forward jump to backwards call), as opposed to just a call?



# Not Relying on the Absolute Address

If we need a string, and we don't know where the shellcode is loaded, we could build this string by moving byte values into a writable location in memory.

But where can we write?

The stack is a very good option. It is writable, and we know its location (ESP).

EIP →	push dword 0x000A2165	ESP →	...
EIP →	push dword 0x646F636C	ESP →	65 21 0A 00
EIP →	push dword 0x6C656853	ESP →	6C 63 6F 64
EIP →	mov eax, esp	EAX ESP →	53 68 65 6C
EIP →	push dword ...	ESP →	...
EIP →	push dword ...	ESP →	...
EIP →	nop	ESP →	...

EAX points to the string "Shellcode!\n" somewhere in memory.

# Filters.

Every exploit has its own constraints on the shellcode. A common one is a set of prohibited byte values (`\0`, `\n`, `\r`, ...). This is called a "filter."

CISC is your friend.

## Setting a register to zero

B8 00 00 00 00    mov eax, 0

31 C0              xor eax, eax

29 C0              sub eax, eax

B8 FF FF FF FF    mov eax, -1  
40                  inc eax

89 D8              mov eax, ebx  
31 D8              xor eax, ebx

## Setting a register to some arbitrary value

```
B8 E4 58 41 00  mov eax, 0x004158E4
```

```
B8 E5 59 40 01  mov eax, 0x014059E5
```

```
35 01 01 01 01  xor eax, 0x01010101
```

```
B8 40 8E 15 04  mov eax, 0x04158E40
```

```
C1 C8 04          ror eax, 4
```

```
B8 1B A7 BE FF  mov eax, 0xFFBEA71B
```

```
F7 D0          not eax
```

Setting a register to some value less than  $2^8$

```
31 C0  xor eax, eax
```

```
B0 FF  mov al, 0xFF
```

Setting a register to some value less than  $2^7$

```
6A 41  push byte 0x41
```

```
58      pop eax
```

## Adding an arbitrary value to a register

```
05 22 11 00 00  add eax, 0x1122
```

```
2D DE EE FF FF  sub eax, -0x1122
```



## Pro Tips.

### Setting a register to zero

Is it even?

```
C1 E0 1F shl eax, 31
```

Is it positive?

```
C1 E8 1F shr eax, 31
```

or

```
C1 F8 1F sar eax, 31
```

## Setting EDX to zero when EAX is zero

```
89 C2  mov  edx,  eax
```

```
99      cdq
```

## Setting EAX, ECX and EDX to zero

```
31 C9  xor ecx, ecx  
F7 E1  mul ecx
```

ECX = 0

EDX:EAX = EAX \* ECX

## Call to unaligned self...

EIP → ...

EIP → E8

FF

FF

FF

EIP → FF

EAX → C0

EIP → 58

EIP → ...

ESP → | ... |  
ESP → | 7A DA D5 AD |

E8 FF FF FF FF

FF C0

58

call -1


inc eax (almost nop)

pop eax

# Encoders.

If it is not feasible to write the shellcode within the given constraints, an encoded shellcode might be necessary.

The decoder must still be written within those constraints, but that is easier than writing all of the shellcode like that.



## Alphanumeric Encoded Shellcode

```
LLLLYhb0pLX5b0pLHSSPPWQPPaPWSUTBRDJfh5tDSRajYX0Dka0Tka  
fhN9fYf1Lkb0TkdfY0Lkf0Tkgh6rfYf1Lki0tkkh95h8Y1LkmjpY  
0Lkq0tkrh2wnuX1Dks0tkwjfX0Dkx0tkx0tkyCjnY0LkzC0TkzCCjt  
X0DkzC0tkzCj3X0Dkz0TkzC0tkzChjG3IY1LkzCCCC0tkzChpfcMX1  
DkzCCCC0tkzCh4pCnY1Lkz1TkzCCCCfhJGfXf1Dkzf1tkzCCjHX0Dk  
zCCCCjvY0LkzCCCjdX0DkzC0TkzCjWX0Dkz0TkzCjdX0DkzCjXY0Lk  
z0tkzMdgvvn9F1r8F55h8pG9wnuvjrNfrVx2LGkG3IDpfcM2KgmnJG  
gbinYshdvD9d
```

# Debugging



INT3 (0xcc) is your friend. It's a software breakpoint.

## shellcode.asm

```
[bits 32]
    int3
    mov eax, 0x11223344
    nop
```

<b>\$ nasm shellcode.asm</b>	<i>assemble your code</i>
<b>\$ ndisasm -u shellcode</b>	<i>check for NUL bytes, etc.</i>
<b>\$ gdb ./demo</b>	<i>launch debugger</i>
<b>\$ gdb --args ./demo shellcode</b>	<i>...with argument</i>

## Sane gdb settings.

Put this in `~/.gdbinit`:

```
set disassembly-flavor intel
set disable-randomization off
set pagination off
set history filename ~/.gdbhistory
set history save
set history expansion
```

(gdb) **run ./shellcode**

Program received signal SIGTRAP, Trace/breakpoint trap.

0xf7fdd001 in ?? ()

(gdb) **x/3i \$eip-1**

0xf7fdd000: int3

=> 0xf7fdd001: mov eax,0x11223344

0xf7fdd006: nop

(gdb) **info registers eax**

eax 0x26718662 644974178

(gdb) **stepi**

0xf7fdd006 in ?? ()

(gdb) **x/3i \$eip-6**

0xf7fdd000: int3

0xf7fdd001: mov eax,0x11223344

=> 0xf7fdd006: nop

(gdb) **info registers eax**

eax 0x11223344 287454020

# Failed System Calls

shellcode.asm

```
; execve syscall  
push byte 11  
pop eax  
int 0x80  
  
; breakpoint  
int3
```

(gdb) **run ./shellcode**

Program received signal SIGTRAP, Trace/breakpoint trap.

0xf7fdd010 in ?? ()

(gdb) **info registers eax**

eax 0xffffffff -2

If EAX is negative, an error occurred. You can find a table of error codes in `/usr/include/asm-generic/errno-base.h`.

```
...  
#define EPERM      1    /* Operation not permitted */  
#define ENOENT     2    /* No such file or directory */  
#define ESRCH      3    /* No such process */  
#define EINTR      4    /* Interrupted system call */  
...
```

```
$ strace ./demo ./shellcode
```

```
...
```

```
execve("|\\215\\25", [0x10], [/* 8 vars */])
```

```
= -1 ENOENT (No such file or directory)
```

Note that the `execve` *library call* returns -1 on error, while the `execve` *system call* returns an error code.

The wrapper in the C library saves the error code in `errno` and returns -1. The `strace` output reflects this, even if you are not using the C library.



# Determining if the shellcode is running

If you can't attach a debugger (e.g., because you are trying to exploit a program on a machine you don't have access to), it can be very valuable to know whether your code is running or not. Usually it is possible to differentiate an infinite loop from a crash.

# Question

What is the shortest byte sequence that results in an infinite loop on x86?

```
eb fe    jmp short -2
```

# Summary

- ▶ Shellcode is a sequence of instructions to be executed in the target process.
- ▶ Shellcode must sometimes adhere to filters (size, allowed bytes, ...).
- ▶ Use `int 0x80` to make system calls.
- ▶ Student presentations: There will be an assignment, where you hand-in your preferred topic(s) and preferences for dates. Priority will be given to uniqueness of topic and flexibility (and extra credit for May 13).