

Proactive Computer Security

Software Reverse Engineering

Morten Brøns-Pedersen
<f@ntast.dk>

(Slides mostly by Morten Shearman Kirkegaard <moki@fabletech.com>)

2016-05-09

The art of taking things apart.

Why?

- To make compatible software
- To see how something works
- To find vulnerabilities

- Executable and Linkable Format (ELF) — Linux, BSD, MINIX, Solaris, IRIX, HP-UX, BeOS, GNU/Hurd, PlayStation, Wii, and a *lot* of embedded systems.
- Portable Executable (PE) — Windows
- Mach-O — Mach, OS X, iOS

What are we dealing with?

```
$ file something.bin
something.bin: ELF 32-bit LSB executable, Intel 80386,
version 1 (SYSV), dynamically linked (uses shared libs),
for GNU/Linux 2.6.26,
BuildID[sha1]=0x1fdc7c28e708e92a297e67448d8a9592bd12b02a,
not stripped
```

```
$ readelf -a something.bin |less
```

```
ELF Header:
```

```
  Magic:   7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
```

```
  Class:                             ELF32
```

```
  Data:                               2's complement, little endian
```

```
  Version:                            1 (current)
```

```
  OS/ABI:                             UNIX - System V
```

```
  ABI Version:                        0
```

```
  Type:                               EXEC (Executable file)
```

```
  Machine:                            Intel 80386
```

```
  Version:                            0x1
```

```
  Entry point address:                0x8048320
```

```
...
```

```
$ readelf -a something.bin |less
```

```
...
```

```
Dynamic section at offset 0x598 contains 25 entries:
```

Tag	Type	Name/Value
0x00000001	(NEEDED)	Shared library: [libc.so.6]
0x0000000c	(INIT)	0x80482b8
0x0000000d	(FINI)	0x80484e0
0x00000019	(INIT_ARRAY)	0x804958c
0x0000001b	(INIT_ARRAYSZ)	4 (bytes)
0x0000001a	(FINI_ARRAY)	0x8049590
0x0000001c	(FINI_ARRAYSZ)	4 (bytes)
0x00000004	(HASH)	0x804818c
0x6ffffef5	(GNU_HASH)	0x80481b4
0x00000005	(STRTAB)	0x8048224

```
...
```

```
$ readelf -a something.bin |less
```

```
...
```

```
Symbol table '.dynsym' contains 5 entries:
```

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	00000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	00000000	0	FUNC	GLOBAL	DEFAULT	UND	puts@GLIBC_2.0
2:	00000000	0	NOTYPE	WEAK	DEFAULT	UND	__gmon_start__
3:	00000000	0	FUNC	GLOBAL	DEFAULT	UND	__libc_start_m
4:	080484fc	4	OBJECT	GLOBAL	DEFAULT	16	_IO_stdin_used

```
...
```


Interesting Sections

- .text — Executable code
- .rodata — Read-only data
- .data — Initialized (read-write) data
- .bss — Uninitialized¹ (read-write) data

¹Actually zero-initialized

After linking an executable, the section information can be stripped away.

The *program headers* will still give us the information we need.

```
$ readelf -a something.bin |less
```

```
...
```

Program Headers:

Type	Offset	VirtAddr	FileSiz	MemSiz	Flg	Align
PHDR	0x000034	0x08048034	0x00100	0x00100	R E	0x4
INTERP	0x000134	0x08048134	0x00013	0x00013	R	0x1
[Requesting program interpreter: /lib/ld-linux.so.2]						
LOAD	0x000000	0x08048000	0x0058c	0x0058c	R E	0x1000
LOAD	0x00058c	0x0804958c	0x00120	0x00124	RW	0x1000
DYNAMIC	0x000598	0x08049598	0x000f0	0x000f0	RW	0x4

```
...
```

R = readable, **W** = writable, **E** = executable

PhysAddr (not shown here) is the same as VirtAddr on modern systems.

To make code compiled using different tool chains compatible, we need calling conventions.

Floating point numbers and structures are more complex, so we will ignore them in this lecture.

cdecl

Arguments are pushed to the stack. The last argument is pushed first. The *caller* removes the arguments from the stack, after the call.

On 80386 the result is stored in EAX (32 bit or less)
or EDX:EAX (up to 64 bits).

The registers EAX, ECX, and EDX are *caller saved*, the rest are *callee saved*.

```
void my_func(int a, int b, int c);
```

```
my_func(1, 2, 3);
```

```
push 3
```

```
push 2
```

```
push 1
```

```
call my_func
```

```
add esp, 12
```

Question

Why do we push the last argument first?

Some functions have a variable number of arguments. The `printf()` function is a classic example.

```
int printf(const char *format, ...);  
  
printf("i is %d, j is %d.\n", i, j);
```

Newer revisions of this convention requires the stack to be 16 byte aligned when entering a function. The alignment is ensured by `main()` and called functions maintain it.

0804848f <main>:

```
804848f:      8d 4c 24 04      lea     ecx, [esp+0x4]
8048493:      83 e4 f0         and     esp, 0xffffffff
8048496:      ff 71 fc         push    DWORD PTR [ecx-0x4]
```

```
void my_func(int a, int b, int c);
```

```
my_func(1, 2, 3);
```

```
mov [esp+8], 3
```

```
mov [esp+4], 2
```

```
mov [esp], 1
```

```
call my_func
```

stdcall

The "*Standard*" *Calling Convention* is the calling convention used by Microsoft for Win32.

Much like *cdecl*; arguments are pushed to the stack, the last argument is pushed first.

But: The *callee* removes the arguments from the stack, when returning.

```
void my_func(int a, int b, int c);
```

```
my_func(1, 2, 3);
```

```
push 3
```

```
push 2
```

```
push 1
```

```
call my_func
```

(One Variation of) fastcall

The first two arguments are stored in ECX and EDX, the rest are pushed to the stack. The last argument is pushed first.

The *callee* removes the arguments from the stack, when returning.

```
void my_func(int a, int b, int c);
```

```
my_func(1, 2, 3);
```

```
push 3  
mov edx, 2  
mov ecx, 1  
call my_func
```

Question

Why do we prefer to use registers
instead of the stack?

Dead Listing (Static Analysis)

We look at the code. We do not execute it.

```
$ strings something.bin  
/lib/ld-linux.so.2  
M?Sn  
__gmon_start__  
libc.so.6  
_IO_stdin_used  
puts  
__libc_start_main  
GLIBC_2.0  
PTRh0  
Hello, World!  
;*2$"
```

```
$ strings something.bin  
/lib/ld-linux.so.2      ← loader/dynamic linker  
M?Sn  
__gmon_start__  
libc.so.6  
_IO_stdin_used  
puts  
__libc_start_main  
GLIBC_2.0  
PTRh0  
Hello, World!  
;*2$"
```

```
$ strings something.bin
/lib/ld-linux.so.2      ← loader/dynamic linker
M?Sn
__gmon_start__
libc.so.6               ← the c library
_IO_stdin_used
puts
__libc_start_main
GLIBC_2.0
PTRh0
Hello, World!
;*2$"
```

```
$ strings something.bin
/lib/ld-linux.so.2      ← loader/dynamic linker
M?Sn
__gmon_start__
libc.so.6               ← the c library
_IO_stdin_used
puts                    ← imported symbol
__libc_start_main
GLIBC_2.0
PTRh0
Hello, World!
;*2$"
```

```
$ strings something.bin
/lib/ld-linux.so.2      ← loader/dynamic linker
M?Sn
__gmon_start__
libc.so.6               ← the c library
_IO_stdin_used
puts                    ← imported symbol
__libc_start_main
GLIBC_2.0
PTRh0
Hello, World!           ← interesting string
;*2$"
```

```
$ readelf -x .rodata ./something.bin
```

```
Hex dump of section '.rodata':
```

```
0x080484f8 03000000 01000200 .....  
0x08048500 48656c6c 6f2c2057 Hello, W  
0x08048508 6f726c64 2100      orld!.
```

```
$ readelf -x .data ./something.bin
```

```
Hex dump of section '.data':
```

```
0x080496a4 00000000 00000000 .....
```

```
$ objdump -M intel -d something.bin >something.txt
$ vim emacs something.txt
...
804840c: 55          push    ebp
804840d: 89 e5       mov     ebp,esp
804840f: 83 e4 f... and     esp,0xfffffffff0
8048412: 83 ec 2... sub     esp,0x20
8048415: 83 7d 0... cmp     DWORD PTR [ebp+0x8],0x1
...
```


When writing programs in C, you may get the impression that execution of your program begins at `main()`. It doesn't.

The compiler will add some code (usually called `_start()`) to initialize your environment. You don't have to know much about this process.

The most important thing is the call to `__libc_start_main()`. The first argument to this call is a pointer to `main()`.

If you have symbols, great.

If not, you can still find `main()`, via the call to `__libc_start_main()`.

Unless the code is malicious—or deliberately obfuscated—you can skip ahead, and begin reverse engineering there. The startup code is not interesting.

Remember the ELF header?

```
$ readelf -a something.bin |less
```

ELF Header:

...

Entry point address: 0x8048320

...

```
8048320: 31 ed      xor    ebp,ebp
8048322: 5e         pop    esi
8048323: 89 e1      mov    ecx,esp
8048325: 83 e4 f... and    esp,0xfffffffff0
8048328: 50         push   eax
8048329: 54         push   esp
804832a: 52         push   edx
804832b: 68 70 8... push   0x8048470
8048330: 68 80 8... push   0x8048480
8048335: 51         push   ecx
8048336: 56         push   esi
8048337: 68 0c 8... push   0x804840c <-- THIS MUST BE MAIN
804833c: e8 cf f... call   8048310 <__libc_start_main@plt>
8048341: f4         hlt
```

```
804840c: 55          push ebp
804840d: 89 e5       mov  ebp,esp
804840f: 83 e4 f...  and  esp,0xfffffffff0
8048412: 83 ec 2...  sub  esp,0x20
8048415: 83 7d 0...  cmp  DWORD PTR [ebp+0x8],0x1
8048419: 75 0e       jne  8048429 <main+0x1d>
804841b: c7 04 2...  mov  DWORD PTR [esp],0x8048500
8048422: e8 c9 f...  call 80482f0 <puts@plt>
8048427: eb 32       jmp  804845b <main+0x4f>
8048429: c7 44 2...  mov  DWORD PTR [esp+0x1c],0x0
```

So, we've found `main()`. How do we figure out what the code does?

We'll begin by looking at how common things in C look in assembly code.

Function Prologue

```
push    ebp
mov     ebp,esp
sub     esp,0x10
```

Function Epilogue

```
mov esp, ebp  
pop ebp  
ret
```

or

```
leave    (= mov esp, ebp ; pop ebp)  
ret      (= pop ip)
```


Padding between Functions

```
90                                nop
66 90                            xchg    ax,ax
8d b4 26 00 00 00 00            lea     esi,[esi+eiz*1+0x0]
8d bc 27 00 00 00 00            lea     edi,[edi+eiz*1+0x0]
```

Arguments and Local Variables

Arguments are at EBP *plus* something:

```
mov    eax,DWORD PTR [ebp+0xc]  
mov    edx,DWORD PTR [ebp+0x8]
```

Local variables are at EBP *minus* something:

```
mov    DWORD PTR [ebp-0x4],eax
```

The compiler is free to *not* use EBP, and just use offsets from ESP.

Pen and paper are your friends. Use them to draw stack frames.

Drawing exercise, yay!

Draw the stack layout for this program. Discuss with your peers.

```
unsigned int foo(unsigned int n, int x) {  
    unsigned int a, b;  
    if (n <= 1) {  
        return x;  
    }  
    a = foo(n - 1, x);  
    b = foo(n - 2, x);  
    return a + b + x;  
}
```

```
x = 123;  
if (a == b) {  
    puts("yes");  
}  
y = 456;
```

```
; -----  
    mov [ebp-16], 123  
    mov eax, [ebp-4]  
    mov edx, [ebp-8]  
    cmp eax, edx  
    jnz .l1  
;  
    mov [esp], 0x8048510  
    call 80482f0 <puts@plt>  
;  
.l1:  
    mov [ebp-12], 456  
;  
-----
```

```
x = 123;
while (i != 0) {
    puts("i is still != 0");
    i--;
}
y = 456;
```

```
; -----  
    mov [ebp-16], 123  
    mov ebx, [ebp-4]  
  
; -----  
.10:  
    test ebx, ebx  
    jz .11  
  
; -----  
    mov [esp], 0x8048510  
    call 80482f0 <puts@plt>  
    dec ebx  
    jmp .10  
  
; -----  
.11:  
    mov [ebp-12], 456  
  
; -----
```



```
x = 123;  
do {  
    puts("some text");  
    i--;  
} while (i != 0);  
y = 456;
```

```
; -----  
    mov [ebp-16], 123  
    mov ebx, [ebp-4]  
;  
-----  
.10:  
    mov [esp], 0x8048510  
    call 80482f0 <puts@plt>  
    dec ebx  
    test ebx, ebx  
    jnz .10  
;  
-----  
    mov [ebp-12], 456  
;  
-----
```

```
x = 123;  
for (i=0; i<10; i++) {  
    puts("some text");  
}  
y = 456;
```

```
; -----  
    mov [ebp-16], 123  
    xor ebx, ebx  
;  
-----  
.10:  
    cmp ebx, 10  
    jge l1  
;  
-----  
    mov [esp], 0x8048510  
    call 80482f0 <puts@plt>  
    inc ebx  
    jmp .10  
;  
-----  
.11:  
    mov [ebp-12], 456  
;  
-----
```

(Trick) Question

How can we tell if it's a `for` loop
or a `while` loop?

Question

What does this function do?

```
80483dc: 55          push    ebp
80483dd: 89 e5       mov     ebp,esp
80483df: 83 ec 10    sub     esp,0x10
80483e2: 8b 45 0c    mov     eax,DWORD PTR [ebp+0xc]
80483e5: 8b 55 08    mov     edx,DWORD PTR [ebp+0x8]
80483e8: 01 d0       add     eax,edx
80483ea: 89 45 fc    mov     DWORD PTR [ebp-0x4],eax
80483ed: 8b 45 fc    mov     eax,DWORD PTR [ebp-0x4]
80483f0: c9         leave
80483f1: c3         ret
```

Coffee Break

```
int arr[3] = { 1, 2, 3 };  
printf("%d\n", arr[1]);
```

```
mov    DWORD PTR [esp+0x1c],0x1  
mov    DWORD PTR [esp+0x20],0x2  
mov    DWORD PTR [esp+0x24],0x3
```

```
mov    eax,DWORD PTR [esp+0x20]
```

```
mov    DWORD PTR [esp+0x4],eax  
mov    DWORD PTR [esp],0x8048530  
call   8048300 <printf@plt>
```


(Trick) Question

How can we tell that it's an array?

```
int arr[3] = { 1, 2, 3 };  
my_func(arr);
```

```
mov     DWORD PTR [esp+0x14],0x1  
mov     DWORD PTR [esp+0x18],0x2  
mov     DWORD PTR [esp+0x1c],0x3
```

```
lea     eax,[esp+0x14]
```

```
mov     DWORD PTR [esp],eax  
call    804841c <my_func>
```

```
void my_func(int *arr)
{
    printf("%d\n", arr[2]);
}
```

```
...
mov     eax,DWORD PTR [ebp+0x8]
add     eax,0x8
mov     eax,DWORD PTR [eax]

mov     DWORD PTR [esp+0x4],eax
mov     DWORD PTR [esp],0x8048500
call    8048300 <printf@plt>
...
```

```
struct my_struct {  
    int a;  
    int b;  
    int c;  
};  
struct my_struct x = { 1, 2, 3 };
```

```
mov     DWORD PTR [esp+0x14],0x1  
mov     DWORD PTR [esp+0x18],0x2  
mov     DWORD PTR [esp+0x1c],0x3
```

```
printf("%d\n", x.b);
```

```
mov     eax,DWORD PTR [esp+0x18]
```

```
mov     DWORD PTR [esp+0x4],eax
```

```
mov     DWORD PTR [esp],0x8048510
```

```
call    8048300 <printf@plt>
```

(Trick) Question

How can we tell that it's a structure?

```
struct my_struct x = { 1, 2, 3 };  
my_func(&x);
```

```
mov     DWORD PTR [esp+0x14],0x1  
mov     DWORD PTR [esp+0x18],0x2  
mov     DWORD PTR [esp+0x1c],0x3
```

```
lea     eax,[esp+0x14]
```

```
mov     DWORD PTR [esp],eax  
call    804841c <my_func>
```

```
void my_func(struct my_struct *ms)
{
    printf("%d\n", ms->b);
}
```

```
mov     eax,DWORD PTR [ebp+0x8]
```

```
mov     eax,DWORD PTR [eax+0x4]
```

```
mov     DWORD PTR [esp+0x4],eax
```

```
mov     DWORD PTR [esp],0x8048500
```

```
call    8048300 <printf@plt>
```


Usually—but not always—elements of structures are accessed as `pointer + constant`, and elements of arrays are accessed as `pointer + variable`.

If you execute my code, I control your computer!

Why do it, then?

It *can* be faster because:

- You can make fewer assumptions
- You can test your hypotheses
- It's easier to concentrate on code that matters

You can use everything you know from static analysis, but now you have data.

If you have symbols, gdb can place breakpoints based on function names.

```
$ gdb ./something.bin
```

```
(gdb) break main  
Breakpoint 1 at 0x804840f
```

```
(gdb) run  
Starting program: ./something.bin
```

```
Breakpoint 1, 0x0804840f in main ()  
(gdb)
```

If you have symbols, gdb also knows which instructions belong to a function.

Breakpoint 1, 0x0804840f in main ()

(gdb) disassemble

Dump of assembler code for function main:

```
0x0804840c <+0>:      push    ebp
0x0804840d <+1>:      mov     ebp,esp
=> 0x0804840f <+3>:      and     esp,0xffffffff
0x08048412 <+6>:      sub     esp,0x20
0x08048415 <+9>:      cmp     DWORD PTR [ebp+0x8],0x1
0x08048419 <+13>:     jne     0x8048429 <main+29>
...
0x0804845b <+79>:     mov     eax,0x0
0x08048460 <+84>:     leave
0x08048461 <+85>:     ret
```

End of assembler dump.

If you don't have symbols, gdb has no idea where a function begins or ends.

```
(gdb) break main
```

```
Function "main" not defined.
```

```
Make breakpoint pending on future shared library  
load? (y or [n])
```

But you can still break on symbols imported from libraries.

```
(gdb) break __libc_start_main  
Breakpoint 1 at 0x8048310
```

```
(gdb) run  
Starting program: ./something.bin
```

```
Breakpoint 1, __libc_start_main (main=0x804840c, argc=1,  
    ubp_av=0xfffffd0b4, init=0x8048480, fini=0x8048470,  
    rtld_fini=0xf7fee590, stack_end=0xfffffd0ac)  
    at libc-start.c:96  
96      libc-start.c: No such file or directory.
```


If you have debugging info for your libc, gdb will know about function arguments.

```
Breakpoint 1, __libc_start_main (main=0x804840c, argc=1,  
    ubp_av=0xfffffd0b4, init=0x8048480, fini=0x8048470,  
    rtld_fini=0xf7fee590, stack_end=0xfffffd0ac)  
    at libc-start.c:96  
96      libc-start.c: No such file or directory.
```

```
(gdb) break *main  
Breakpoint 2 at 0x804840c
```

Otherwise we'll just examine the stack. We know that ESP points to the return address, and that the first argument (at ESP+4) is a pointer to main().

```
(gdb) x/2wx $esp  
0xffffd08c:      0x08048341      0x0804840c
```

```
(gdb) break *0x0804840c  
Breakpoint 2 at 0x804840c
```

If we are looking for vulnerabilities, we may care about *how* a program does something. Otherwise we only care about *what* it does.

A good reverse engineer will know what to reverse and what to skip.

What does my_func() do?

```
0x08048339 <+9>:      lea      eax, [esp+0x10]
0x0804833d <+13>:     mov      DWORD PTR [esp], eax
0x08048340 <+16>:     mov      DWORD PTR [esp+0x4], 0x4d2
0x08048348 <+24>:     call     0x8048440 <my_func>
0x0804834d <+29>:     ...
```

```
(gdb) break *0x0804834d  
Breakpoint 1 at 0x804834d
```

```
(gdb) run  
Starting program: ./something.bin
```

```
Breakpoint 1, 0x0804834d in main ()
```

```
(gdb) x/12bx $esp+0x10
```

0xffffc000:	0x31	0x32	0x33	0x34
0xffffc004:	0x00	0xe5	0xfe	0xf7
0xffffc008:	0x8b	0x84	0x04	0x08

```
(gdb) x/s $esp+0x10
```

```
0xffffc000: "1234"
```

Where is a function called from?

```
(gdb) break my_func  
Breakpoint 1 at 0x8048440
```

```
(gdb) run  
Starting program: ./something.bin
```

```
Breakpoint 1, 0x08048440 in my_func ()
```

```
(gdb) backtrace  
#0  0x08048440 in my_func ()  
#1  0x0804834d in main ()
```

Objects are just structures. If they have virtual methods, the first element of the structure is a pointer to an array of function pointers.

It's not easy to do static analysis of C++ programs.

You can learn a lot about how the compiler generates code, by compiling and debugging tiny programs. Do it.

There are plenty of reverse engineering challenges out there. Read other people's writeups, and solve some challenges.

In CTF competitions, they are often labeled *binary challenges*. In other places they are called *crack-mes* and *keygen-mes*. Beware that some can be extremely difficult.

- <http://pwnable.kr>
- <http://treasure.pwnies.dk>

That's it. Have fun!